



Week 10 Workshop - Database Transactions

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Transactions

- A **transaction** is a sequence of database operations grouped together for execution as a logic unit in a DBMS.

Steps	Transaction
	BEGIN TRANSACTION
1	SELECT balance FROM ACCOUNT WHERE name = 'Steve';
2	UPDATE ACCOUNT SET balance = balance-100 WHERE name = 'Steve';
3	SELECT balance FROM ACCOUNT WHERE name = 'Bob';
4	UPDATE ACCOUNT SET balance = balance+500 WHERE name = 'Bob';
5	COMMIT;

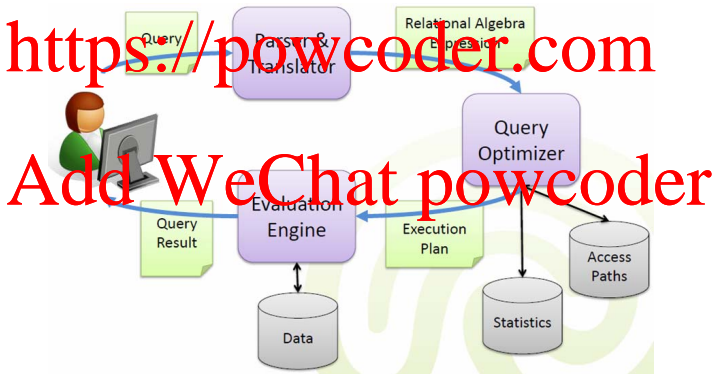
<https://powcoder.com>

Add WeChat powcoder



Transactions

- What's the difference between database transactions and programs written by a programming language like C, Java, and Python?
- How are transactions handled in the query processing?



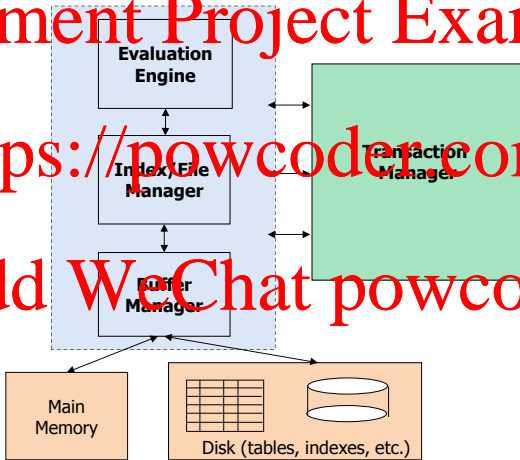


Transaction Manager - A Simplified View

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





Transactions - ACID Properties

Assignment Project Exam Help

Transactions

```
T1 : BEGIN TRANSACTION  
      SELECT ...  
      UPDATE ...  
      COMMIT
```

```
T2 : SELECT ...
```

```
T3 : INSERT ...
```

```
T4 : BEGIN TRANSACTION  
      SELECT ...  
      DELETE ...  
      ABORT
```

ACID properties

Atomicity

Consistency

Isolation

Durability

<https://powcoder.com>


Add WeChat powcoder

Transactions - ACID Properties

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





Transactions - ACID Properties

Assignment Project Exam Help

ACID properties

Atomicity

Consistency

Isolation

Durability

Transaction Manager

Recovery


Concurrency
control

Consistency is the responsibility of an application developer.

Transaction Manager - Common Techniques

Assignment Project Exam Help

<https://powcoder.com>



- **Add WeChat powcoder**
Logging for recovery—assuring atomicity/durability of transactions
e.g., Write-Ahead Log (WAL) Protocol



Logging - Introduction

Assignment Project Exam Help

- A **transaction log** is an append-only file that records changes to objects made by transactions.
- When multiple transactions run concurrently, log records are interleaved.
- A transaction log can be implemented as a separate file or set of files in the database.
- Recovery amounts to **either undoing or redoing changes from the log**:
 - **Undo** the operations that have not been committed;
 - **Redo** the operations that have been committed but not yet been written to disk.
- **Checkpoints** tell the points from which to begin applying transaction logs during database recovery.

<https://powcoder.com>

Add WeChat powcoder



Write-Ahead Log (WAL) Protocol

Assignment Project Exam Help

- Write-Ahead Log (WAL) requires that a record of every change to a database is available while attempting to recover from a crash.

Any change to an object is first recorded in the log, i.e., a record containing both the old and new values for the object.

- A record in the log must be written to persistent storage before committing the transaction.

- Accordingly, the definition of a **committed transaction** is:

“A transaction, all of whose log records, including a commit record, have been written to persistent storage”.



Write-Ahead Log (WAL) Protocol

- Typical fields in a log record:



- Each log record has a unique id called **LSN** (Log Sequence Number).

- prevLSN** is the LSN of the previous log record written by the same transaction.

- Possible **types** include: update, commit, abort, end, etc.

- Does WAL bring in some benefits for performance?

- Often results in a significantly reduced number of disk writes
- Supports one sync against the log file instead of potentially many against the data files
- Enables online backup and point-in-time recovery



Transaction Manager - Recovery

Assignment Project Exam Help

- Key concepts to aid in recovery:

- **Transaction log**: records of database operations

<https://powcoder.com>

Write Ahead Log (WAL)

- **Undo ...**

- **Redo ...**

Add WeChat powcoder


- **Checkpoint**: snapshot of the state of a database

(Widely used in practice, but not covered in this course)

Transaction Manager - Common Techniques

Assignment Project Exam Help

<https://powcoder.com>



- **Logging** for recovery – assuring **atomicity/durability** of transactions
e.g., Write-Ahead Log (WAL) Protocol
- **Locking** for concurrency control – assuring **isolation** of transactions
e.g., Two-Phase Locking (2PL) Protocol



Locking - Introduction

Assignment Project Exam Help

- A **lock** is associated with an object, e.g., file, table, record, page, etc.
- Two main types of locks:
 - **Shared lock** (read-lock): for reading an object by a transaction
 - **Exclusive lock** (write-lock): for writing an object by a transaction

(Note: there are other types of locks defined by different DBMSs)

- Lock compatibility:

Lock type	read-lock	write-lock
read-lock	Yes	No
write-lock	No	No


Two-Phase Locking (2PL) Protocol

Assignment Project Exam Help

- Locks are handled in two phases:
 - **Expanding**: locks are acquired and no locks are released.
 - **Shrinking**: locks are released and no locks are acquired.

<https://powcoder.com>

Add WeChat powcoder



Two-Phase Locking (2PL) Protocol


Assignment Project Exam Help

Bad news:

- 2PL can radically limit interleaving among transactions in some cases ...
- 2PL may be subject to **deadlocks**, i.e., the mutual blocking of two or more transactions

<https://powcoder.com>

Step	T_1	T_2
1	lock-r(A)	
2	read(A)	
3		lock-r(B)
4		read(B)
5	lock-w(B)	
6	write(B)	
7		lock-w(A)
8		write(A)



- T_1 is waiting for T_2 to get a write-lock on B . T_2 is waiting for T_1 to get a write-lock on A .



Two-Phase Locking (2PL) Protocol

Assignment Project Exam Help

Good news:

- 2PL makes interleaving safe, i.e., guarantees the serializability property for transactions.
- **Serializability** means that a resulting database state is equal to a database state of running transactions serially.
- Serializability is the major correctness criterion for concurrent transactions.





Serializability - Example

- Consider $A = 200$ and $B = 500$, and we have two concurrent transactions:



- Serializable transactions:





Serializability - Example



- Consider $A = 200$ and $B = 500$, and we have two concurrent transactions:



- Case 1:** $A=900$ and $B=800$


- Case 2:** $A=700$ and $B=900$

- Serializable transactions:



Serializability - Example


- Consider $A = 200$ and $B = 500$, and we have two concurrent transactions:



- Case 1:** $A=900$ and $B=800$

- Case 2:** $A=700$ and $B=900$

- Are the following transactions serializable?




- Yes. $A=900$ and $B=800 \leftrightarrow$ equivalent to Case 1!



Serializability - Example

- Consider $A = 200$ and $B = 500$, and we have two concurrent transactions:



- Case 1:** $A=900$ and $B=800$

- Case 2:** $A=700$ and $B=900$

- Are the following transactions serializable?



- No. $A=900$ and $B=900 \leftrightarrow$ not equivalent to Case 1 or Case 2!



Problems in Concurrent Transactions

Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:

<https://powcoder.com>

**Lost
update**

write \rightarrow write

**Dirty
read**

write \rightarrow read

**Unrepeatable
read**

read \rightarrow write
(read)

**Phantom
read**

read \rightarrow write
(read)

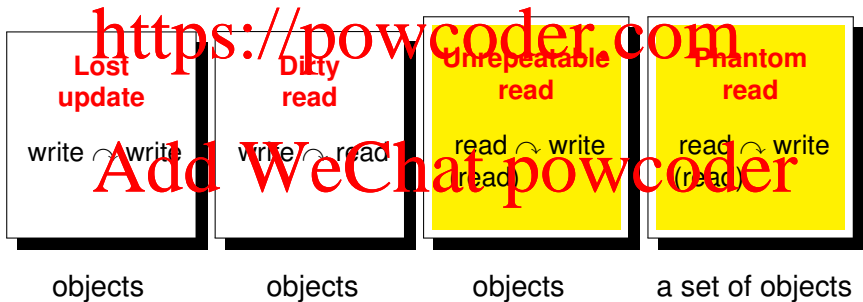
Add WeChat powcoder



Problems in Concurrent Transactions

Assignment Project Exam Help

- If no concurrency control for transactions, some problems may occur:





The Lost Update Problem - Another Example

- Ben and Amy have the same salary. T_1 sets their salaries to \$80,000, and T_2 sets their salaries to \$90,000.

- If executing T_1 and T_2 sequentially,
 - for T_1 ; T_2 , both receive \$90,000.

<https://powcoder.com>
 → Either is acceptable from the transaction viewpoint.

- If executing T_1 and T_2 concurrently, we may have:

	T_1	T_2
1	write(A) (A:=80000)	
2		write(A) (A:=90000)
3		write(B) (B:=90000)
4		commit
5	write(B) (B:=80000)	
6	commit	

→ It is not acceptable!



The Dirty Read Problem - Another Example

- Both Ben and Amy are rewarded a bonus \$5,000 and a pay rise 5%. T_1 increases her salaries with \$5,000 and T_2 increments their salaries by 5%.

- If executing T_1 and T_2 sequentially, they would have Also, T_1 or T_2 could abort for some reasons. \hookrightarrow all are acceptable from the transaction viewpoint.

<https://powcoder.com>

- If executing T_1 and T_2 concurrently, we may have:

	T_1	T_2
1	read(A)	
2	write(A) ($A := A + 5000$)	
3		read(A)
4	read(B)	
5	write(B) ($B := B + 5000$)	
6	abort	
7		write(A) ($A := A + A \times 5\%$)
8		read(B)
9		write(B) ($B := B + B \times 5\%$)
10		commit

\hookrightarrow It is not acceptable!

The Unrepeatable Read Problem - Another Example

Assignment Project Exam Help

- Amy and Ben are using a website to book flight tickets to Brisbane.
 - Amy signs on first to see that only one ticket is left, and finds it expensive.
 - Amy takes time to decide. Ben signs on later and also finds one ticket left, orders it instantly, and logs off.
 - Amy decides to buy a ticket, and finds no tickets left.

<https://powcoder.com>

Add WeChat powcoder

	T_1 (from Amy)	T_2 (from Ben)
1	read(X)	
2		read(X)
3		write(X) ($X := X - 1$)
4		commit
5	read(X)	

- This situation can never arise in a serial execution of T_1 and T_2 .



The Phantom Read Problem - Another Example

- Amy is 30 years old, but her age in the table players is mistakenly recorded as 40. Ben is 25 years old and his age is correctly recorded in players.
- Suppose that we have the following two current transactions:

```
T1: SELECT * FROM players
    WHERE age<32;
    .
    SELECT * FROM players;
    WHERE age<32;
    COMMIT;
```

```
T2: UPDATE players
    SET age=30
    WHERE rating=8 and name='Amy';
    COMMIT;
```

Add WeChat powcoder

	T ₁	T ₂
1	read(players)	
2		read(players)
3		write(players)
4		commit
5	read(players)	
6	commit	

- This situation also can never arise in a serial execution of T_1 and T_2 .



Discussion

Assignment Project Exam Help

What are the differences between 'unrepeatable read' and 'phantom read'?

- **Unrepeatable read**

- Executing the same SELECT twice yields the same tuples, but attribute values might be different;
- May occur when reading objects that are affected by UPDATE from another transaction;
- Can be prevented using record-level locking.

- **Phantom read**

- Executing the same SELECT twice yields two different sets of tuples;
- May occur when querying a set of tuples that are affected by INSERT/DELETE/UPDATE from another transaction;
- Can be prevented using table-level locking.



What Should We lock?

- Consider the following two concurrent transactions again:

```
T1: SELECT * FROM players
    WHERE age<32;
...
SELECT * FROM players
    WHERE age<32;
COMMIT;
```

```
T2: UPDATE players
    SET age=30
    WHERE rating=8 and name='Amy';
COMMIT;
```

- What objects should the DBMS lock in order to avoid the phantom read problem?

- Table-level locks**

- e.g., read-lock on players for T_1 , write-lock on players for T_2

- Record-level locks**

- e.g., read-lock on every record with age<32 for T_1 , write-lock on every record with rating=8 and name='Amy' for T_2

- ...

Transaction Support in SQL

- An explicit transaction may have no **BEGIN TRANSACTION** statement but must be ended with either **COMMIT** or **ABORT (ROLLBACK)** statement.

- When no explicit transaction statements are given, each single SQL statement is considered to be a transaction.
- To give programmers more control over transaction overhead, SQL allows them to specify **isolation level**, i.e., the degree of interference that a transaction is prepared to tolerate on concurrent transactions.

- Key idea:

To trade off **consistency** (i.e., increased risk of violating database integrity) with **performance** (i.e., greater concurrent access to data)



Isolation Levels

Assignment Project Exam Help

- SQL-92 defines four isolation levels:

1 Read Uncommitted

2 Read Committed

3 Repeatable Reads

4 Serializable

- To specify an isolation level, e.g.,

```
SET TRANSACTION ISOLATION LEVEL serializable;
```

- The SQL standard does not impose a specific locking scheme or mandate particular behaviors.

<https://powcoder.com>

Add WeChat powcoder



Isolation Levels

Assignment Project Exam Help

- The intention is to prohibit certain problems:

Isolation Level	Dirty Read	Unrepeatable Read	Phantom Read
READ UNCOMMITTED	Yes	Yes	Yes
READ COMMITTED	No	Yes	Yes
REPEATABLE READ	No	No	Yes
SERIALIZABLE	No	No	No

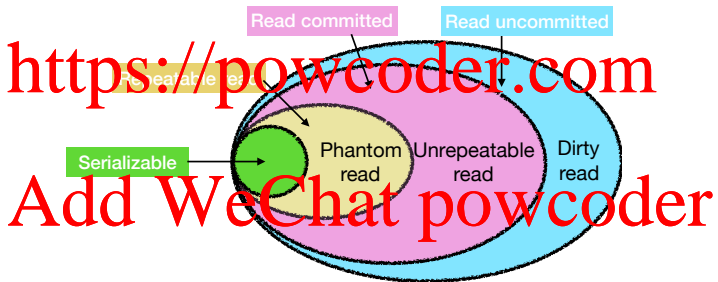
- Different DBMSs implement isolation levels quite differently.
- The isolation level required for **Lost Update** is debatable (depending on a DBMS's implementations). But in general, it may require the highest level **SERIALIZABLE** to prevent it.¹

¹ https://drtom.ch/posts/2011/11/12/The_Lost_Update_Problem_-_Part_1/

Isolation Levels - Concurrency Control

Assignment Project Exam Help

- A DBMS provides different levels of isolation → different degrees of concurrency control to prevent different problems.




- Concurrency control is **NOT** binary in a database system.




Isolation Levels - Read Uncommitted

- Read Uncommitted is the least restrictive isolation level.
- One transaction can see changes made by other transactions which are not yet committed. This can be quite dangerous.
- Use it when executing queries over read-only data or if it does not matter whether a query returns uncommitted data.



Isolation Levels - Read Committed

- **Read Committed**: One transaction only sees committed changes by other transactions.
- It is the most commonly used isolation level in database applications.
- Use it when you want to maximize concurrency between applications but do not want queries to see uncommitted data.



Isolation Levels - Repeatable Reads

Assignment Project Exam Help

- **Repeatable Reads:** The objects touched by a transaction are locked and cannot be updated or deleted by a concurrent transaction.
- Use it when you want some level of concurrency between applications but do not expect individual objects to be changed during a transaction.

<https://powcoder.com>

Add WeChat powcoder



Isolation Levels - Serializable


Assignment Project Exam Help

- Serializable is the highest isolation level. All transactions are totally isolated from other transactions. It is safe but may cause significant performance hit.

- Use it when you want some level of concurrency between applications but do not expect that a query returns different sets of results when running at different times.

<https://powcoder.com>

Add WeChat powcoder




Locks Taken by SQL Server for Isolation Levels ²

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



² <http://michaeljswart.com/2012/06/visualizing-transaction-isolations-for-sql-server/>



Wrap-up - Isolation Levels

Assignment Project Exam Help

- A lower isolation level increases the ability of many users to access data **at the same time**, but also increases the number of concurrency effects (such as dirty reads or lost updates) users might encounter.
- Conversely, a higher isolation level **reduces the types of concurrency effects that users may encounter**, but requires more system resources and increases the chances that one transaction will block another.
- Choosing the appropriate isolation level depends on **balancing**
 - **the data integrity requirements of the application**
against
 - **the overhead of each isolation level.**



Assignment Project Exam Help

<https://powcoder.com>

Research Topics

Add WeChat powcoder



Research Topics

Assignment Project Exam Help

- This is an active research area covering many interesting research topics.
- Historically, much of the work has been done in the context of relational database systems.
- However, the ideas in general are independent of whether the underlying system is a relational database system or something else.

<https://powcoder.com>

Add WeChat powcoder

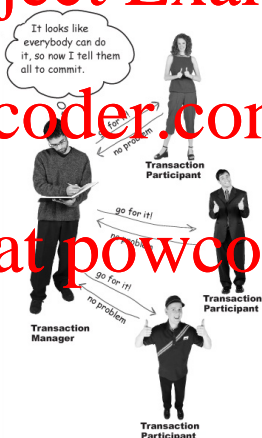
- Distributed database systems
- Graph database systems
- Document-oriented database systems
- ...

Research Topics

• Distributed transactions

Phase ONE

Phase TWO



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder